

```
In [1]: # Import Libraries
import numpy as np
import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler
from warnings import filterwarnings
filterwarnings('ignore')
```

```
In [3]: data = pd.read_csv('C:/Users/saile/Desktop/IoT/datasets/model_evaluation_spine_dataset.csv')
```

```
In [4]: data.head()
```

```
Out[4]:
```

	Col1	Col2	Col3	Col4	Col5	Col6	Col7	Col8	Col9	Col10
0	63.027817	22.552586	39.609117	40.475232	98.672917	-0.254400	0.744503	12.5661	14.5386	15.3046
1	39.056951	10.060991	25.015378	28.995960	114.405425	4.564259	0.415186	12.8874	17.5323	16.7848
2	68.832021	22.218482	50.092194	46.613539	105.985135	-3.530317	0.474889	26.8343	17.4861	16.6589
3	69.297008	24.652878	44.311238	44.644130	101.868495	11.211523	0.369345	23.5603	12.7074	11.4246
4	49.712859	9.652075	28.317406	40.060784	108.168725	7.918501	0.543360	35.4940	15.9546	8.8723

```
In [5]: data.shape
```

```
Out[5]: (310, 14)
```

```
In [6]: # Data is clean except the "Unnamed: 13" column
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 310 entries, 0 to 309
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Col1        310 non-null    float64
1   Col2        310 non-null    float64
```

```

2  Col3          310 non-null    float64
3  Col4          310 non-null    float64
4  Col5          310 non-null    float64
5  Col6          310 non-null    float64
6  Col7          310 non-null    float64
7  Col8          310 non-null    float64
8  Col9          310 non-null    float64
9  Col10         310 non-null    float64
10 Col11         310 non-null    float64
11 Col12         310 non-null    float64
12 Class_att     310 non-null    object
13 Unnamed: 13   14 non-null    object
dtypes: float64(12), object(2)
memory usage: 34.0+ KB

```

```

In [7]: # Type of Backbone Conditions
        data.Class_att.unique()

```

```

Out[7]: array(['Abnormal', 'Normal'], dtype=object)

```

```

In [8]: # Remove unwanted column
        df = data.drop("Unnamed: 13", axis=1)

```

```

In [9]: # Change the Column names to be sensible
        df.rename(columns = {
            "Col1" : "pelvic_incidence",
            "Col2" : "pelvic_tilt",
            "Col3" : "lumbar_lordosis_angle",
            "Col4" : "sacral_slope",
            "Col5" : "pelvic_radius",
            "Col6" : "degree_spondylolisthesis",
            "Col7" : "pelvic_slope",
            "Col8" : "direct_tilt",
            "Col9" : "thoracic_slope",
            "Col10" : "cervical_tilt",
            "Col11" : "sacrum_angle",
            "Col12" : "scoliosis_slope",
            "Class_att" : "target"}, inplace=True)

```

```


In [10]: # How skewed is the data?
          df["target"].value_counts().sort_index().plot.bar()

```

```

Out[10]: <AxesSubplot:>

```



```
In [11]: # Convert categorical to numeric {"Abnormal":0, Normal:1}
df.target = df.target.astype("category").cat.codes
```

```
In [12]: df.head()
```

```
Out[12]:
```

	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	degree_spondylolisthesis
0	63.027817	22.552586	39.609117	40.475232	98.672917	-0.254400
1	39.056951	10.060991	25.015378	28.995960	114.405425	4.564259
2	68.832021	22.218482	50.092194	46.613539	105.985135	-3.530317
3	69.297008	24.652878	44.311238	44.644130	101.868495	11.211523
4	49.712859	9.652075	28.317406	40.060784	108.168725	7.918501

```
In [13]: # 88% Accuracy
dataset = df[["pelvic_incidence", "pelvic_tilt", "lumbar_lordosis_angle", "sacral_slope", "
```

```
In [14]: # Separate input and output
y = dataset.target
X = dataset.drop("target", axis=1)
```

```
In [15]: # Split data between train and test
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=
```

```
In [16]: # List models
from sklearn import datasets
from sklearn.cluster import KMeans
models = [LogisticRegression, DecisionTreeClassifier, RandomForestClassifier, GradientB
```

```
In [17]: # Train & Predict models
acc_list = []
name_list = []
for model in models:
    clf = model()
    clf = clf.fit(X_train, y_train)
    predictions = clf.predict(X_test)
    name_list.append((model).__name__)
    acc_list.append(classification_report(y_test, predictions, output_dict=True)["accuracy"])
    print((model).__name__, "--> ", classification_report(y_test, predictions, output_dict=True)
```

```
LogisticRegression --> 0.8640776699029126
DecisionTreeClassifier --> 0.7378640776699029
RandomForestClassifier --> 0.8252427184466019
GradientBoostingClassifier --> 0.7961165048543689
SVC --> 0.7961165048543689
```

```
In [18]: # Make a dataframe
team = pd.DataFrame(list(zip(name_list, acc_list)))
team.columns = ['Name', 'Accuracy']
team
```


```
Out[18]:
```

	Name	Accuracy
0	LogisticRegression	0.864078
1	DecisionTreeClassifier	0.737864
2	RandomForestClassifier	0.825243
3	GradientBoostingClassifier	0.796117
4	SVC	0.796117

```
In [19]: # Render a Chart
sns.barplot(x=team["Name"], y=team["Accuracy"], data=team)

# Rotate x-Labels
plt.xticks(rotation=-45)
plt.ylim(0.7, 1)
```

```
Out[19]: (0.7, 1.0)
```



```
In [20]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
```

```
In [21]: km = KMeans(n_clusters=2, random_state=42)
#
# Fit the KMeans model
#
km.fit_predict(X_train)
#
# Calculate Silhouette Score
#
score = silhouette_score(X_train, km.labels_, metric='euclidean')
#
# Print the score
#
print('Silhouetter Score: %.3f' % score)
```

Silhouetter Score: 0.478

```
In [22]: pip install yellowbrick
```

Requirement already satisfied: yellowbrick in c:\users\saille\anaconda3\lib\site-packages (1.5)
Requirement already satisfied: scikit-learn>=1.0.0 in c:\users\saille\anaconda3\lib\site-packages (from yellowbrick) (1.1.3)
Requirement already satisfied: matplotlib!=3.0.0,>=2.0.2 in c:\users\saille\anaconda3\lib\site-packages (from yellowbrick) (3.4.3)
Requirement already satisfied: scipy>=1.0.0 in c:\users\saille\anaconda3\lib\site-packages (from yellowbrick) (1.7.1)
Requirement already satisfied: numpy>=1.16.0 in c:\users\saille\anaconda3\lib\site-packages (from yellowbrick) (1.20.3)
Requirement already satisfied: cyclor>=0.10.0 in c:\users\saille\anaconda3\lib\site-packages (from yellowbrick) (0.10.0)
Requirement already satisfied: six in c:\users\saille\anaconda3\lib\site-packages (from c

```

ycluster>=0.10.0->yellowbrick) (1.16.0)
Requirement already satisfied: pillow>=6.2.0 in c:\users\saille\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (8.4.0)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\saille\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\saille\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.3.1)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\saille\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (3.0.4)
Requirement already satisfied: joblib>=1.0.0 in c:\users\saille\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->yellowbrick) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\saille\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->yellowbrick) (2.2.0)
Note: you may need to restart the kernel to use updated packages.

```

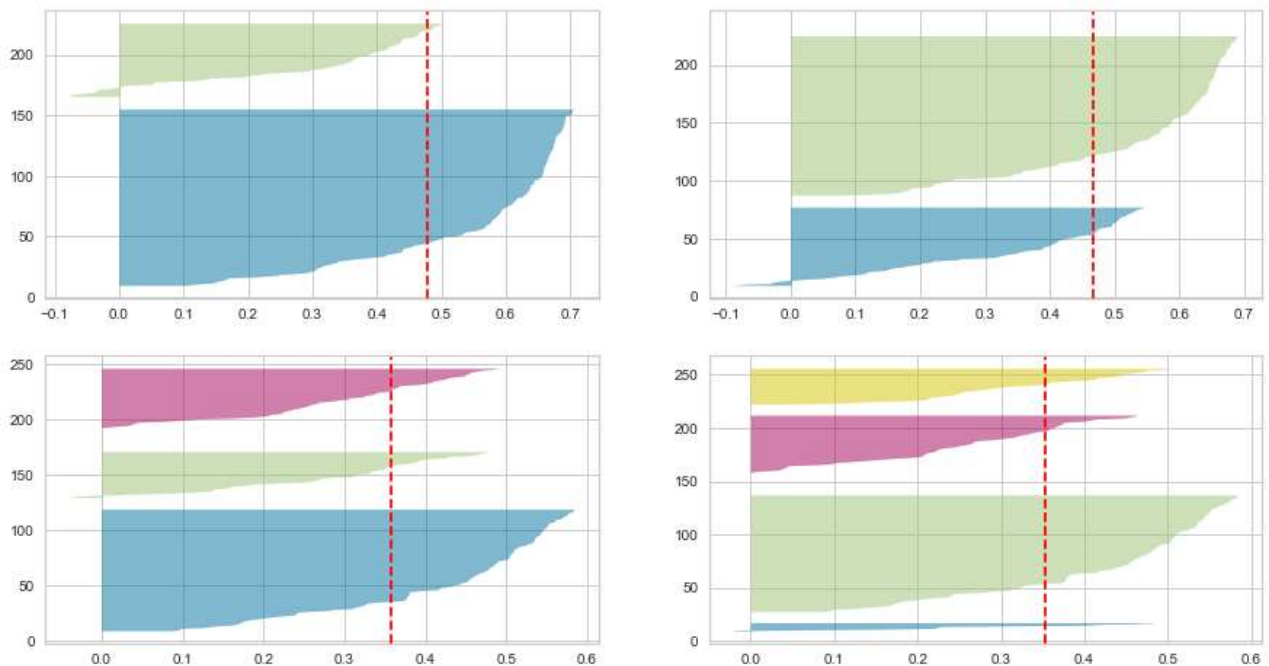
In [23]:

```

from yellowbrick.cluster import SilhouetteVisualizer

fig, ax = plt.subplots(2, 2, figsize=(15,8))
for i in [2,3,4,5]:
    '''
    Create KMeans instance for different number of clusters
    '''
    km = KMeans(n_clusters=i,random_state=42)
    q, mod = divmod(i, 2)
    '''
    # Create SilhouetteVisualizer instance with KMeans instance
    #Fit the visualizer
    '''
    visualizer = SilhouetteVisualizer(km, colors='yellowbrick', ax=ax[q-1][mod])
    visualizer.fit(X_train)

```



In []: